

## Problem Statement

### **Title: Boston Housing Price Prediction using Linear Regression via Deep Neural Network**

#### **Objective:**

To predict median house prices in Boston suburbs using the Boston Housing dataset by implementing **linear regression** using a **deep neural network (DNN)** approach with TensorFlow/Keras.

---

## Dataset Description:

The **Boston Housing Dataset** includes 506 samples and 13 numerical/categorical features such as:

- CRIM: Crime rate per capita
- RM: Average number of rooms per dwelling
- LSTAT: % lower status of the population
- PTRATIO: Pupil-teacher ratio, etc.

**Target Variable:** MEDV (Median value of owner-occupied homes in \$1000s)

---

## Methodology:

Although this is a **linear regression** problem, we implement it using a **deep neural network** configured to learn a linear function:

#### **Steps:**

1. **Import Libraries** (TensorFlow, Keras, NumPy, Pandas, etc.)
2. **Load the Dataset** (from sklearn.datasets or keras.datasets)
3. **Preprocess the Data**
  - Normalize features
  - Train/test split
4. **Define the Neural Network Architecture**
  - Input layer matching number of features
  - Only 1 dense layer with **no activation** (linear regression)
  - Loss function: **Mean Squared Error**

- Optimizer: **SGD** or **Adam**

## 5. **Train the Model**

## 6. **Evaluate the Model** on test set

## 7. **Visualize** the predicted vs actual house prices

## ✓ **Expected Outcome**

A trained DNN model that approximates linear regression to predict house prices, evaluated using Mean Absolute Error (MAE) or Mean Squared Error (MSE).

## ✦ **Goal of the Practical**

You are using the **Boston Housing dataset** to predict house prices (regression problem) using a **neural network model built with TensorFlow/Keras**.

---

## 🔍 **Step-by-step Explanation**

---

### ◆ **1. Importing Required Libraries**

```
import tensorflow as tf

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

import numpy as np

import matplotlib.pyplot as plt
```

### ✓ **Explanation:**

- tensorflow: Used to build and train the neural network.
  - sklearn: For splitting data and normalization.
  - matplotlib.pyplot: To visualize predictions.
  - numpy: Numerical operations.
-

## ◆ 2. Load the Boston Housing Dataset

```
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.boston_housing.load_data()
```

### ✓ Explanation:

- Loads the dataset from Keras.
  - `x_train, x_test`: Feature values (13 input variables like crime rate, room count, etc.).
  - `y_train, y_test`: Target variable (house price in \$1000s).
- 

## ◆ 3. Normalize the Features

```
scaler = StandardScaler()
```

```
x_train_scaled = scaler.fit_transform(x_train)
```

```
x_test_scaled = scaler.transform(x_test)
```

### ✓ Explanation:

- `StandardScaler`: Scales features to have **mean = 0** and **std = 1**.
  - Normalizing helps the model learn faster and improves accuracy.
- 

## ◆ 4. Build the DNN Model (Linear Regression)

```
model = tf.keras.models.Sequential([  
    tf.keras.layers.Input(shape=(13,), name='input-layer'),  
    tf.keras.layers.Dense(100, name='hidden-layer-2'),  
    tf.keras.layers.BatchNormalization(name='hidden-layer-3'),  
    tf.keras.layers.Dense(50, name='hidden-layer-4'),  
    tf.keras.layers.Dense(1, name='output-layer')  
])
```

### ✓ Explanation:

- `Sequential`: Stacks layers in order.
- `Input(shape=(13,))`: Specifies that input has 13 features.
- `Dense(100)`: Fully connected hidden layer with 100 neurons.
- `BatchNormalization`: Stabilizes and speeds up training.

- Dense(50): Another hidden layer with 50 neurons.
- Dense(1): Output layer with **1 neuron** (for predicting a single number = house price).

💡 Even though it's called a DNN, this is used for **regression**, not classification.

---

## ◆ 5. Compile the Model

```
model.compile(optimizer='adam', loss='mse', metrics=['mae'])
```

### ✅ Explanation:

- optimizer='adam': Adaptive optimizer (good default choice).
  - loss='mse': Mean Squared Error (used for regression).
  - metrics=['mae']: Mean Absolute Error shown during training.
- 

## ◆ 6. Train the Model

```
model.fit(x_train_scaled, y_train, epochs=100, validation_split=0.2)
```

### ✅ Explanation:

- Trains the model for 100 epochs.
  - Uses 20% of training data for validation.
  - Adjusts model weights to minimize error.
- 

## ◆ 7. Evaluate the Model

```
loss, mae = model.evaluate(x_test_scaled, y_test)
```

```
print(f"Test MAE: {mae}")
```

### ✅ Explanation:

- Evaluates model performance on **unseen test data**.
  - MAE (Mean Absolute Error) tells us how far predictions are from actual values on average.
- 

## ◆ 8. Make and Plot Predictions

```
y_pred = model.predict(x_test_scaled)
```

```
plt.scatter(y_test, y_pred)

plt.xlabel("Actual Prices")

plt.ylabel("Predicted Prices")

plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], '--r')

plt.show()
```

#### ✅ Explanation:

- `model.predict()`: Makes predictions on test data.
- scatter plot: Compares actual vs predicted prices.
- Red line: Ideal line where prediction = actual. Closer points are to the line → better model.

#### ✅ Summary for Viva or Written:

- You used a DNN to perform linear regression.
- The dataset had 13 features; the model had input → hidden layers → output.
- Used MSE loss function and Adam optimizer.
- Normalized data before training.
- Evaluated with MAE and visualized predictions.

#### 🔍 What the columns mean:

Layer Type	Output Shape	Description
<b>InputLayer</b>	(None, 13)	Accepts input of 13 features
<b>Dense (100)</b>	(None, 100)	100 neurons, connected to all 13 inputs
<b>BatchNormalization</b>	(None, 100)	Normalizes 100 outputs
<b>Dense (50)</b>	(None, 50)	50 neurons
<b>Dense (1)</b>	(None, 1)	Final output: 1 value (price prediction)